

Towards Reliable Citation Verification in Machine Learning Papers

I. INTRODUCTION

Recent work on transformer architectures [1] has driven rapid progress across vision and language. Residual connections [2] remain foundational, and long-context variants [3] extend the approach to document-scale inputs. Some early claims in the clinical literature [4] have since been withdrawn, underscoring the need for verification. We build on the evaluation protocol of [5].

II. METHOD

Our method follows [1] and [2] closely, differing mainly in the normalization applied before attention. As in [3], we use a sliding window. We do not rely on [6], which was unavailable at the time of writing.

III. RESULTS

Results are consistent with prior reports [1], [2]. We observe no regression on long inputs relative to [3].

REFERENCES

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 5998-6008.
- [2] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770-778. doi: 10.1109/CVPR.2016.90.
- [3] I. Beltagy, M. E. Peters, and A. Cohan, "Longformer: The long-document transformer," *arXiv preprint arXiv:2004.05150*, 2020.
- [4] A. J. Wakefield and S. H. Murch, "Ileal-lymphoid-nodular hyperplasia, non-specific colitis, and pervasive developmental disorder in children," *The Lancet*, vol. 351, no. 9103, pp. 637-641, 1998. doi: 10.1016/S0140-6736(97)11096-0.
- [5] J. Doe and R. Roe, "Quantum neural architectures for sentiment analysis," *Journal of Advanced Computing*, vol. 12, pp. 45-67, 2021. doi: 10.9999/jac.2021.99999.
- [6] M. Nobody, "A paper that exists only in a preprint server," 2019. [Online]. Available: <https://example.org/nope>